

Spam Detection Model Report

1. Introduction

With the growing number of digital communications, spam messages have become a significant problem. This project aims to develop a machine learning model that can accurately classify SMS messages as either "spam" or "ham" (non-spam). The solution involves natural language processing (NLP) techniques and a Naive Bayes classification model to detect spam messages effectively.

2. Dataset

Source:

The dataset used is the SMS Spam Collection, a well-known corpus containing 5,572 labeled messages collected for spam detection research.

Structure:

- v1: Label of the message (ham or spam)
- v2: Message content (raw text)

Preprocessing:

- Unnecessary columns were dropped.
- Columns were renamed for clarity: label (spam or ham) and text (the actual message).

3. Text Preprocessing

To prepare the text data for modeling, a cleaning function was implemented. The steps include:

- Lowercasing: Standardizing text by converting all characters to lowercase.
- Punctuation Removal: Stripping all punctuation to reduce noise.
- Stopword Removal: Eliminating common English words that do not contribute to classification (e.g., "the", "and", "is").
- Stemming: Reducing words to their base or root form using the Porter Stemmer.

A new column, `clean_text`, was added to the dataset containing the cleaned version of each message.

4. Feature Extraction

To convert text data into numerical format for machine learning:

- TF-IDF Vectorization was applied.

- Maximum of 3000 features were used to limit dimensionality and reduce overfitting.
- The resulting output is a feature matrix representing each message numerically.

5. Model Building

Algorithm Used:

- Multinomial Naive Bayes: A simple yet effective algorithm suitable for text classification problems.

Process:

- The dataset was split into 80% training and 20% testing using `train_test_split`.
- The Naive Bayes model was trained on the training data.

6. Evaluation Metrics

The model was evaluated using the following metrics:

- Accuracy: ~98%
- Precision: High
- Recall: High
- F1-Score: High

Additional outputs included:

- Confusion Matrix: Showing true positives, true negatives, false positives, and false negatives.
- Classification Report: Detailed per-class performance.

7. Prediction Function

A utility function `predict_message(msg)` was implemented to:

- Clean and preprocess a new message.
- Vectorize the message using the trained TF-IDF vectorizer.
- Predict and return the label (Spam or Ham).

Example Predictions:

"Congratulations! You won a free iPhone. Click here to claim." → Spam

"Hey, are we still on for lunch today?" → Ham


8. Model Persistence

The trained model and vectorizer were saved using Python's pickle module:

- `spam_model.pkl`: Contains the trained Naive Bayes model.
- `vectorizer.pkl`: Contains the TF-IDF vectorizer.

These files can be reused in a web app or any production environment for real-time spam detection.

9. UI using Streamlit



Spam Message Detector


Enter your message:

CHECK SPAM

Prediction Result

This message is: NOT Spam

Confidence: 0.91

 **Safety Tips:**

- Do not click unknown links or download attachments from suspicious messages.
- Never share personal, banking, or login details via text.
- Spam often uses urgency, prizes, or fear tactics to trick you. Stay alert!

Check Another Message

10. Conclusion

This project successfully demonstrates how machine learning and NLP can be combined to develop a robust spam detection system. The model shows high accuracy and is efficient for deployment. Future improvements may include:

- Testing with more diverse datasets.
- Exploring deep learning approaches.
- Enhancing the front-end user interface.